

# System State Forecast

**Predicting User State via Resource Metrics**

January 31, 2026

System Forecast Team

# Contents



1	Project Overview .....	3
1.a	Executive Summary .....	4
1.b	Phase 1: The Pivot .....	5
2	System Architecture .....	6
2.a	Three Main Stages .....	7
2.b	Data Collection .....	8
3	The Engineering .....	9
3.a	Feature Engineering .....	10
3.b	The “Idle” Detection Challenge .....	11
4	Model Performance .....	12
4.a	Results .....	13
4.b	Visualizing State .....	14
5	Conclusion .....	15
5.a	File Structure .....	16
5.b	Future Improvements .....	17

# 1 Project Overview

# Executive Summary



**Goal:** Classify computer state into **Idle**, **Interactive (Light)**, or **Media Watching** in real-time.

**Constraint:** Non-invasive monitoring. No key/mouse logging, only system resources (CPU, RAM, Disk, Network, GPU).

## Outcome:

- Random Forest Classifier with 83% accuracy.
- Solved the “Idle vs. Reading” distinction using hardware-aware heuristics (GPU Sleep States).

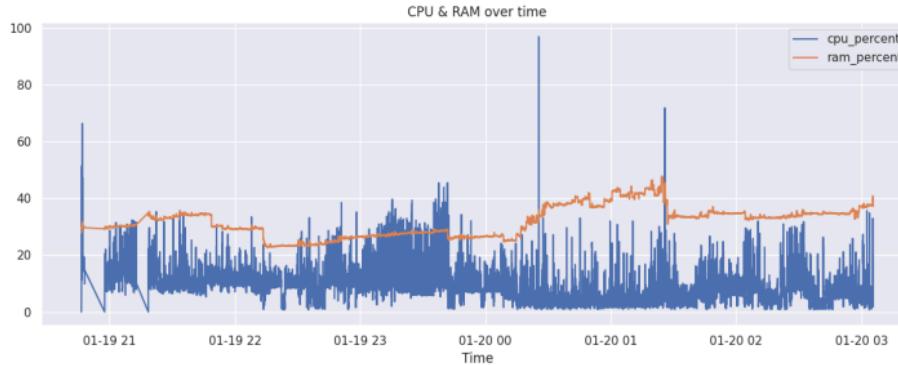
# Phase 1: The Pivot



**Initial Goal:** Predict CPU/RAM usage to forecast heavy loads.

## The Problem:

- **Skewed Distribution:** System is idle/low-usage 99% of the time.
- **Unpredictable Peaks:** Spikes are user-driven, not history-driven.

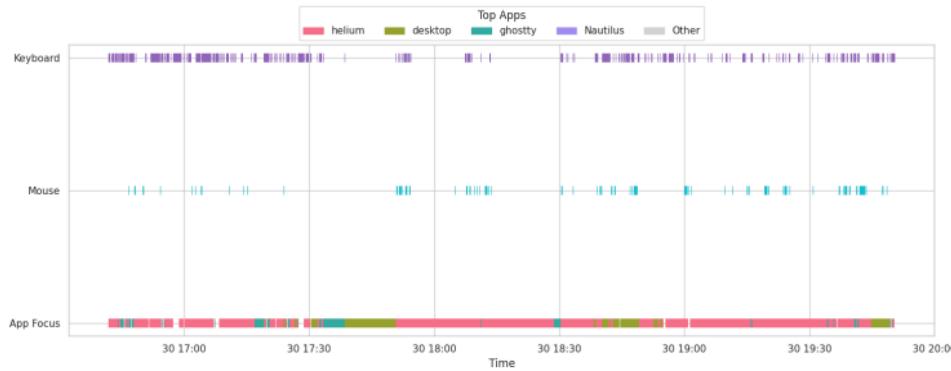


## **2 System Architecture**

# Three Main Stages



- 1. Data Collection:** Background daemon (`10_comprehensive_activity_log.py`) sampling at 1.0s.
- 2. Feature Engineering:** Rolling windows (5s & 30s) to capture trends and variance.
- 3. Live Inference:** Random Forest Model + Heuristic Overrides.



# Data Collection



## Metrics:

- `psutil`: CPU, RAM, Disk I/O, Network.
- `intel_gpu_top`: RC6 (Sleep), RCS (3D), VCS (Video).
- **Ground Truth**: `libinput` (used only for training labels).

**Verification:** We used a `nu` script to visualize data flow in real-time notifications.

```
media_watching ⓘ
notify-send
{cpu_percent: 4.9, max_gpu: 2.213414, idle_time_sec: 1.0, app_id: com.mitchellh.ghostty}
nu watch_metrics.nu cpu_percent max_gpu idle_time_sec app_id
or: nu::shell::error
operation interrupted
-[ /home/pampam/Documents/UTH/Do_an_data/all_of_the_info_necessary/watch_
cs.nu:19:5 ]
```

### 3 The Engineering

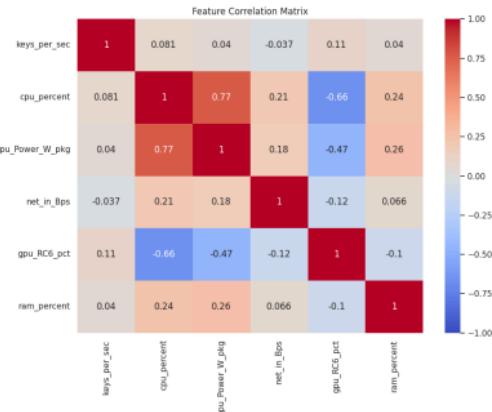
# Feature Engineering



Raw metrics are too volatile. We calculate **Rolling Features**:

- **Short Window (5s)**: Immediate changes (e.g., “Start typing”).
- **Long Window (30s)**: Sustained trends (e.g., “Watching Movie”).

**Key Insight:** Standard Deviation (`std`) is crucial. “Idle” has low variance; “Reading” has micro-bursts.



# The “Idle” Detection Challenge



## Problem:

- **Idle:** CPU 1-2%, GPU Sleep 99%.
- **Reading:** CPU 2-3%, GPU Sleep 97%.
- Statistical distributions overlap > 80%.

## Solution: `max_gpu` & Heuristics

- **Feature:** `max_gpu` (highest load on **any** engine).
- **Heuristic Override:**

```
if prediction == 'interactive_light':  
    if (rc6 > 99.0 and cpu < 5.0) or (max_gpu < 1.0):  
        prediction = 'Idle'
```

## **4 Model Performance**

# Results



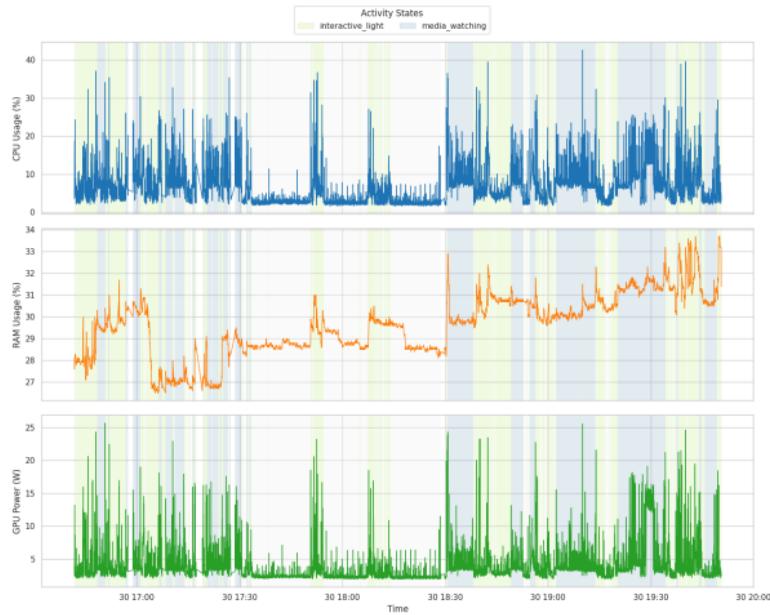
**Algorithm:** Random Forest Classifier

## Performance:

- Accuracy: 83-87%
- High Precision for “Media Watching” (97%).
- “Idle” requires the heuristic override for reliability.

Class	Precision	Recall	F1-Score
Idle	0.70	1.00	0.82
Interactive	0.83	0.97	0.89
Media	0.97	0.77	0.86

# Visualizing State



## 5 Conclusion

# File Structure



- `10_comprehensive_activity_log.py`: Data Collector.
- `system_only_model/`:
  - `train_model.py`: Trainer.
  - `live_inference.py`: Production Inference.
  - `activity_model.joblib`: Trained Model.

# Future Improvements



- **Hysteresis:** Add a state latch (e.g., “Must be Idle for 3s to switch”) to prevent flickering.
- **Per-Process Analysis:** Hard-coded rules for known media players (mpv, vlc).
- **Profiles:** Different baselines for “Gamer” vs “Coder”.

**Thank You**